

# A Step-by-step Tutorial to Reproduce *MonEx* Use Case Experiments

This file contains the experimental artifacts of the use case experiments of the *MonEx* paper. In the same folder of this file, you can find the experiments folders that contain the datasets and some dependent scripts.

## 1 Disk & Power Usage of a *MongoDB* Cluster

### 1.1 Objective

Using *MonEx* framework to monitor the disk & the power usage of a *MongoDB* cluster while running an indexation workload.

### 1.2 Requirements

The experiment has an automated script to run on the *Grid'5000* testbed, but we list the necessary steps to allow repeating it elsewhere.

#### 1.2.1 Hardware Requirements

This experiment needs a cluster of at least five nodes. Three nodes serve as shard nodes while another node holds the configuration server and the *mongo* router daemons. The last node is reserved for installing the *MonEx* framework. Additionally, the main point during reproducing this experiment is on the way we monitor the experiment using *MonEx* but not its results, so any hardware specifications could be used. Thus, the only impact is that you might obtain a different execution time.

#### 1.2.2 Software Requirements

Many software are necessary:

- *MonEx* configured with *Prometheus* as a data collector.
- *Prometheus* default node exporter
- *MongoDB* version 3.2
- *WiredTiger* search engine
- *Ruby* (1.9 or above)
- *mongo*(v1.1), *yaml* ruby-gems to generate data for *MongoDB*

### 1.3 Installation

For simplicity reasons, we refers to the experiment nodes as follows:

- node\_[1-3]: the first, second & third shard nodes of *MongoDB*
- node\_4: the node that holds the config. server and the router of *MongoDB*
- node\_monex: the node that will be used to hold the *MonEx* framework.

#### 1.3.1 *MongoDB* Installation & Configuration

⇒ Please refer to the *MongoDB* official webpage to install *MongoDB* on the nodes\_[1-4] and make sure to stop the *mongod* service on all nodes after the installation:

<https://docs.mongodb.com/v3.2/tutorial/install-mongodb-on-debian/>

⇒ Configure the *WiredTiger* to be used as the default storage engine in *MongoDB*. This is done by editing the file `/etc/mongod.conf` on the nodes\_[1-4], by adding the following line:

```
storage.engine: wiredTiger
```

- ⇒ Configure the *MongoDB* cluster as the following:  
 - run the following commands on the node\_4:

```
nohup mongod --configsvr --port 27019 & // to start a config serv. daemon
nohup mongos --configdb [node_4_IP:27019] // to start a mongo router daemon
```

- run the following command on the nodes\_[1-3]:

```
nohup mongod --config /etc/mongod.conf --shardsvr --port 27017 &
```

- ⇒ We need to configure the cluster and to have a sharded data collection, so we run the following commands on the node\_4 using *mongo* daemon:

```
$ mongos
mongos> sh.addShard([node_1:IP]) // repeat this step for node_2 & node_3
mongos> use db1 // create db1 database and jump inside
mongos> sh.enableSharding(db1) // indicating that the db1 could be distributed
mongos> db.c1.ensureIndex({_id: hashed}) // create a free collection c1
mongos> sh.shardCollection(db1.c1, {_id:hashed}) // sharding the free collection c1
```

- ⇒ Create and inject a data into the configured cluster. To do so, we need to use the provided file (indexing\_on\_MongoDB/generateDumpsDataForMongoDB.rb) for generating data (20 million documents by default, about 80Gbyte). The script could be run on a standalone *MongoDB* and then we could create a dump of the generated data in order to be injected into our cluster using the *mongo* router. The commands on this steps are as follows:

```
ruby generateDumpsDataForMongoDB.rb // on a machine running a mongod daemon
mongodump --db db1 --collection c1 // export the collection data into a bson file
```

On the node\_4, we execute the following command to inject the created data into our cluster

```
mongorestore -d db1 -c c1 [bson file] // inject and evenly distribute the data
```

At the end of this step, *MongoDB* is installed and configured correctly.

### 1.3.2 MonEx installation

Clone the *MonEx* github project into the node\_monex and install its dependencies as explained in the project page (<https://github.com/madynes/monex>). This step implies installing the default exporter of *Prometheus* into all the data nodes (nodes\_[1-3]), and configuring *Prometheus* to regularly scrape those data nodes at the rate of 1 second.

## 1.4 Experiment workflow

We suppose that *MongoDB* is installed correctly and that *MonEx* server is started as explained in the previous steps. Then, we execute the following steps.

1. Before running the main workload of the experiment, notify *MonEx* that the indexation workload is to be started in a moment. Thus, the following command must be run on the node\_4:

```
curl -X POST MonEx_IP:5000/exp/mongoXP
```

2. We then run the workload from the same node as follows:

```
$mongos
mongos> use db1
mongos> db.c1.createIndex({randNum:1})
```

This will take several minutes regarding the hardware specification of the shard nodes, and the size of the generated data.

3. When the workload is accomplished, we should notify *MonEx* that the experiment is terminated. So, the following command is run dynamically from the experiment script (you can run it manually when you the previous command is terminated):

```
curl -X PUT MonEx_IP:5000/exp/mongoXP
```

4. We can query the *MonEx* server to obtain a dataset containing the target metrics of the experiment. The commands to obtain the disk usage, and the power usage metrics are as follows:

```
curl -X GET -H "Content-Type: application/json" IP:5000/exp/mongoXP \
-d '{"query":"rate(node_disk_io_time_ms{device=\"sda1\"}, \
"server\":\"prometheus\",\"type\":\"duration\",\"labels\":[\"instance\"]}' > dataset_hdd.csv
```

```
curl -X GET -H "Content-Type: application/json" IP:5000/exp/mongoXP \
-d '{"query":"pdu","server\":\"prometheus\",\"type\":\"duration\", \
"labels\":[\"outlet\",\"instance\"]}' > dataset_pdu.csv
```

5. The obtained datasets from the previous step are available in the experiment folder. So if it is not possible to repeat the experiment, you can use the *MonEx-draw* over the obtained datasets to produce the figures of this experiment as follows:

```
./monex-draw -F2 dataset_pdu.csv -F dataset_hdd.csv
-c node1,node2,node3
-c2 node1,node2,node3
-x "Time (sec)" -y "Disk utilization (%)" -y2 "Power usage (W)" -l
```

This command produces the *Fig. 2-b* that is appeared in the *MonEx* paper.

## 2 Many-nodes Bittorrent Download

### 2.1 Objective

We want to monitor the completion of 100 peers on a torrent network using *MonEx*.

### 2.2 Requirement

This experiment relies on *Distem* which is easy to deploy on Grid'5000, but it should be repeatable on any testbed.

### 2.3 Installation

To setup this experiment, we need to install *Distem* on 11 machine. On Grid'5000, it is done as follows:

```
frontend> oarsub -t deploy -l slash_22=1+{"cluster='grisou'}nodes=11,walltime=3 -I
frontend> kadeploy3 -f $OAR_NODE_FILE -e jessie-x64-nfs -k
frontend> distem-bootstrap
```

On the coordinator, we deploy the topology using the provided ruby script:

```
coordinator> ruby deploy_peers.rb
```

This script deploy 101 virtual node running Transmission (100 leechers, 1 seeder) and a tracker. It also throttle the bandwidth between peers (you can check details on *Distem* website: <http://distem.gforge.inria.fr/>). We will install *Prometheus* and *MonEx* directly on the coordinator.

```
coordinator> apt install python3-pip python3-flask python3-requests r-base prometheus
coordinator> pip3 install flask-cors
coordinator> systemctl stop prometheus
coordinator> git clone https://github.com/madynes/monex.git
```

Start *Prometheus* and *MonEx* with the provided configuration files

```
coordinator> prometheus -config.file prometheus-config.yml
coordinator> ./monex-server monex_server.conf
```

The *prometheus* configuration file contains the list of peers and the *MonEx* configuration file defines the target (*prometheus*) and a metric. Specifying the metric here is not necessary but it is more convenient.

## 2.4 Experiment workflow

We are now ready to start the experiment, the provided script will start the experiment and signal MonEx

```
./start_xp.sh 127.0.0.1:5000
```

This script make the file available to the seeder and start a leecher every 4 seconds and start an exporter on every virtual node (exporter.py). At the end of the experiment, to get the metric, we use :

```
curl -X GET -H "Content-Type: application/json" IP:5000/exp/peerxp \
-d '{"query":"percent","server":"prometheus", \
"type":"duration", "labels":["node"]}' > dataset.csv
```

But since we have specify the metric in the configuration file of MonEx, we can juste use:

```
curl -X GET -H "Content-Type: application/json" IP:5000/exp/peerxp \
-d '{"query":"completion","type":"duration"} > dataset.csv
```

Moreover, monex-draw can connect directly to the server to plot a figure:

```
./monex-draw -S 127.0.0.1:5000 --metric "completion" --exp "peer" --type "sample" \
-y "Completion" -x "Time (sec)" -n
```

## 3 Input/output offset sequences experiment

### 3.1 Objective

Using *MonEx* framework to monitor how the access patterns over a given file are uncovered when applying a random access workload. An *eBPF* program is used to uncover the access patterns, while *Fio benchmark* is used to generate the random access workload. The *eBPF* program is out of scope of this tutorial as its paper is not yet submitted. Thus, it is only possible to repeat the analysis of this experiment.

### 3.2 Requirements

#### 3.2.1 Hardware Requirements

The experiment must run on any Linux distribution with a kernel version greeter than v3.4 to allow running *eBPF* virtual machine. Additionally, no special hardware is needed.

#### 3.2.2 Software Requirements

- Fio Benchmark v2.1.3
- eBPF virtual machine
- *MonEx* configured with *InfluxDB* as a data collector

### 3.3 Installations

- Fio can be installed using *apt* on debian:

```
apt-get install -y fio
```

- *eBPF* virtual machine could be installed as a part of the *IOvisor BCC project*:  
<https://github.com/iovisor/bcc> (see install.md inside the project page).

### 3.4 Experiment workflow

The experiment is performed as follows:

1. We create a file using *dd* command as follows:

```
dd if=/dev/zero of=testFile bs=16 count=9175040 // create 147 MByte file
```

2. We start our *eBPF* program that listens inside the kernel to new I/O requests accessing the target file
3. Notify MonEx that the experiment will begin in a moment:

```
\begin{verbatim} curl -X PUT MonEx_IP:5000/exp/accessPattern
```

4. Start the main workload of accessing *testFile* randomly:

```
 fio --name=testFile --iodepth=16 --rw=read --direct=0 --numjobs=1 --group_reporting
```

5. After the termination of the workload, notify *MonEx* that the experiment is finished.

```
 curl -X PUT MonEx_IP:5000/exp/accessPattern
```

6. *MonEx* could be then quired to obtain a dataset of the target metrics as follows:

```
 curl -X GET 127.1:5000/exp/myxp -H "Content-Type: application/json" -d \
 '{"measurement":"ebpf","server":"influx","database":"monex","type":"sample"}' > dataset.csv
```

7. Use *MonEx-draw* to produce a publishable figure of the experiment, by consuming the provided dataset file:

```
 ./monex-draw -F dataset.csv -x "Time (sec)" -y "Offset (bytes)" -p -r "random" -n -s ';' ;'
```

This command is used to produce the *Fig. 4* of the *MonEx* paper.